AARON M. MUELLER

CONTACT

Center for Language & Speech Processing
Johns Hopkins University
Hackerman 319, 3400 N. Charles St.
Baltimore, MD 21218-2608 (U.S.A.)

amueller@jhu.edu
aaronmueller.github.io
github.com/aaronmueller

RESEARCH INTERESTS

- Natural language processing
- Robust generalization, interpretability, evaluation
- Multilinguality
- Computational psycholinguistics, syntax, morphology

EDUCATION

Johns Hopkins University — Baltimore, MD
Ph.D., Computer Science — May 2023 (expected)
M.S.E., Computer Science — May 2020
GPA: 3.9/4.0
Advisors: Tal Linzen, Mark Dredze

New York University — New York, NY
Visiting academic, Center for Data Science — Aug. 2021 – May 2023
Advisor: Tal Linzen

University of Kentucky — Lexington, KY
B.S., Computer Science. *Honors* — May 2018
B.S., Linguistics. *Honors* — May 2018
GPA: 4.0/4.0. *Summa cum laude*

EXPERIENCE

Meta — Menlo Park, CA
Research Intern, AI Integrity — May – Nov. 2022
Manager: Kanika Narang
- Research in retrieval-augmented generative models for few-shot question answering.
- Resulted in improved F_1 on multiple QA and classification datasets using far fewer parameters than state-of-the-art models. Also resulted in a submission to ACL (under review) [2].

Amazon Web Services (AWS) — Santa Clara, CA
Applied Scientist Intern, Lex — May – Aug. 2021
Manager: Saab Mansour
- Research in pre-training methods for improving goal-oriented dialogue agents.
- Resulted in state-of-the-art few-shot intent classification accuracy (>30% 1-shot gains) and a publication at ACL [6].

Raytheon BBN Technologies — Cambridge, MA
Research Intern, Analytics & Machine Intelligence — May – Aug. 2019
Manager: Ilana Heintz
- Research in low-resource cross-lingual word alignment and entity linking.
- Implemented convolutional neural machine translation models rivaling our prior seq2seq model's BLEU with over 20% faster training and over 50% faster inference.

PUBLICATIONS

Peer-reviewed Proceedings & Articles

1. **Aaron Mueller**, Tal Linzen. "How to Plant Trees in Language Models: Data and Architectural Effects on the Emergence of Syntactic Inductive Biases." In *Association for Computational Linguistics (ACL)*, 2023.

2. **Aaron Mueller**, Kanika Narang, Lambert Mathias, Qifan Wang, Hamed Firooz. "Meta-learning with Demonstration Retrieval for Efficient Few-shot Learning." In *Findings of the Association for Computational Linguistics (ACL)*, 2023.

3. Koustuv Sinha, Jon Gauthier, **Aaron Mueller**, Kanishka Misra, Keren Fuentes, Roger Levy, Adina Williams. "Language Model Acceptability Judgements Are Not Always Robust to Context." In *Association for Computational Linguistics (ACL)*, 2023.

4. Julian Michael, Ari Holtzman, Alicia Parrish, **Aaron Mueller**, Alex Wang, Angelica Chen, Divyam Madaan, Nikita Nangia, Richard Yuanzhe Pang, Jason Phang, Samuel R. Bowman. "What Do NLP Researchers Believe? Results of the NLP Community Metasurvey." In *Association for Computational Linguistics (ACL)*, 2023.

5. **Aaron Mueller**, Robert Frank, Tal Linzen, Luheng Wang, Sebastian Schuster. "Coloring the Blank Slate: Pre-training Imparts a Hierarchical Inductive Bias to Sequence-to-sequence Models." In *Findings of the Association for Computational Linguistics (ACL)*, 2022.

6. **Aaron Mueller**, Jason Krone, Salvatore Romeo, Saab Mansour, Elman Mansimov, Yi Zhang, Dan Roth. "Label Semantic Aware Pre-training for Few-shot Text Classification." In *Proceedings of the Association for Computational Linguistics (ACL)*, 2022.

7. **Aaron Mueller**, Yu Xia, Tal Linzen. "Causal Analysis of Syntactic Agreement Neurons in Multilingual Language Models." In *Proceedings of the Conference on Computational Natural Language Learning (CoNLL)*, 2022.

8. Alexandra DeLucia, Shijie Wu, **Aaron Mueller**, Carlos Aguirre, Mark Dredze, Philip Resnik. "BERNICE: A Multilingual Pre-trained Encoder for Twitter." In *Proceedings of the Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2022.

9. **Aaron Mueller**, Mark Dredze. "Fine-tuning Encoders for Improved Monolingual and Zero-shot Polylingual Neural Topic Modeling." In *Proceedings of the North American Chapter of the Association for Computational Linguistics (NAACL)*, 2021.

10. **Aaron Mueller**, Zach Wood-Doughty, Silvio Amir, Mark Dredze, Alicia L. Nobles. "Demographic Representation and Collective Storytelling in the Me Too Twitter Hashtag Activism Movement." In *Proceedings of the Association for Computing Machinery (ACM) on Human-Computer Interaction (HCI), vol. CSCW1*, 2021.

11. Matthew Finlayson*, **Aaron Mueller***, Sebastian Gehrmann, Stuart Shieber, Tal Linzen, Yonatan Belinkov. "Causal Analysis of Syntactic Agreement Mechanisms in Neural Language Models." In *Proceedings of the Association for Computational Linguistics (ACL)*, 2021. [*Equal contribution]

12. Alexandra DeLucia*, **Aaron Mueller***, Xiang Lisa Li, João Sedoc. "Decoding Methods for Neural Narrative Generation." In *Proceedings of the Workshop on Generation Evaluation and Metrics (GEM) at Association for Computational Linguistics (ACL)*, 2021. [*Equal contribution]

13. **Aaron Mueller**, Garrett Nicolai, Panayiota Petrou-Zeniou, Natalia Talmina, Tal Linzen. "Cross-linguistic Syntactic Evaluation of Word Prediction Models." In *Proceedings of the Association for Computational Linguistics (ACL)*, 2020.

14. **Aaron Mueller**, Garrett Nicolai, Arya D. McCarthy, Dylan Lewis, Winston Wu, David Yarowsky. "An Analysis of Massively Multilingual Neural Machine Translation for Low-Resource Languages." In *Proceedings of the Language Resources and Evaluation Conference (LREC)*, 2020.

15. Arya D. McCarthy, Rachel Wicks, Dylan Lewis, **Aaron Mueller**, Winston Wu, Oliver Adams, Garrett Nicolai, Matt Post, David Yarowsky. "The Johns Hopkins University Bible Corpus: 1600+ Tongues for Typological Exploration." In *Proceedings of the Language Resources and Evaluation Conference (LREC)*, 2020.

16. Garrett Nicolai, Dylan Lewis, Arya D. McCarthy, **Aaron Mueller**, Winston Wu, David Yarowsky. "Fine-grained Morphosyntactic Analysis and Generation Tools for More Than One Thousand Languages." In *Proceedings of the Language Resources and Evaluation Conference (LREC)*, 2020.

17. Marten van Schijndel, **Aaron Mueller**, Tal Linzen. "Quantity Doesn't Buy Quality Syntax with Neural Language Models." In *Proceedings of the Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2019.

18. Arya D. McCarthy, Winston Wu, **Aaron Mueller**, Bill Watson, David Yarowsky. "Modeling Color Terminology Across Thousands of Languages." In *Proceedings of the Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2019.

19. **Aaron Mueller***, Yash Kumar Lal*. "Sentence-Level Adaptation for Low-Resource Neural Machine Translation." In *Proceedings of the Workshop on Technologies for Machine Translation of Low-Resource Languages (LoResMT) at Machine Translation Summit (MTSummit)*, 2019. [*Equal contribution]

INVITED TALKS

Planting Trees in Language Models: Emergent Syntactic Behaviors and Mechanisms from Pre-training.
- Bar-Ilan NLP Seminar, Bar-Ilan University (Ramat Gan, Israel). Dec. 13, 2022.
- NLP Seminar, Technion – Israel Institute of Technology (Haifa, Israel). Dec. 14, 2022.

What Generalizations do Sequence-to-sequence Models Learn from Multilingual Text? Insights from Translation and Syntactic Transformations. Multilingual Text Processing Group, National Research Council of Canada (Ottawa, ON). Mar. 4, 2022.

Syntactic Agreement in Neural Language Models: How Well and Where Do They Perform Subject-Verb Agreement? Language & Understanding Group, Mila – Québec Artificial Intelligence Institute (Montréal, QC). Mar. 22, 2021.

Causal Mediation Analysis for Analyzing Neural Networks. Fairness & Interpretability Research Talk Series, Google (New York, NY). Mar. 17, 2021.

Causal Analysis of Syntactic Agreement Mechanisms in Neural Language Models. Center for Language & Speech Processing Seminar, Johns Hopkins University (Baltimore, MD). Feb. 12, 2021.

FELLOWSHIPS AND AWARDS

Microsoft Accelerate Foundation Models Research Award, *International* 2023
Awarded for research on the capabilities of large language models. Provides OpenAI API credits and priority GPT-4 access. ($5,000)

National Science Foundation Graduate Research Fellow, *National* 2018 - 2023

Gaines Fellow, *University of Kentucky* 2016 - 2018
Two-year fellowship. Requires the completion of a juried project, a thesis project, and a seminar in the humanities. ($5,000)

Patterson Scholar, *University of Kentucky* 2014 - 2018
Four-year scholarship covering tuition, educational materials, and room & board. Awarded to undergraduates who have earned National Merit semifinalist standing or higher. ($86,000)

Goldwater Scholarship (Honorable Mention), *National* 2017

Phi Beta Kappa, *National* 2017

Raymond F. Betts Scholar, *University of Kentucky* 2017
Awarded for thesis research. Used funds to design language technologies for low-resource dialects of French. ($2,500)

Linguistics Research Award, *University of Kentucky* 2016
Awarded to an undergraduate to facilitate a year-long research project in linguistics. ($500)

TEACHING

Johns Hopkins University Baltimore, MD
Teaching Assistant
- Machine Learning: AI System Design & Development. Spring 2020

MENTORING

Master's students:

- Dan Pechi (New York University). In progress. Work on imparting inductive biases to language models. 2023
- Swapnil Sharma (New York University). In progress. Work on evaluating summarization models. 2022–2023
- Yash Kumar Lal (Johns Hopkins). Met biweekly. Resulted in a workshop publication [19]. 2018–2019

Undergraduate researchers:

- Yu Xia (New York University). Met weekly. Resulted in a publication at CoNLL [7]. 2021–2022
- Matthew Finlayson (Harvard). Met weekly. Resulted in a publication at ACL [11]. 2020–2021

SERVICE

Organizing Committees:

- The BabyLM Shared Task (CoNLL 2023)
- The Inverse Scaling Prize (2022)

Reviewing:

- ACL Rolling Review (Oct. 2021 – present; monthly)
- ACL (2022, 2020)
- EMNLP (2022, 2019)
- CoNLL (2022)
- TACL (2022)
- NAACL (2021)
- CSCW (2021)
- COLING (2020)

SKILLS

Programming:

- Languages: Python, C++, HTML, CSS, Javascript, Bash
- Machine Learning Toolkits: PyTorch (incl. HuggingFace, fairseq, sockeye), NLTK, Scikit-learn, numpy
- Version Control: DVCS (Git, Bitbucket)

Linguistic Tools:

- Praat, AntConc, QGIS, Audacity

LANGUAGES

English (native language), French (B2, Canadian).
Experience in German and Finnish through research.