

# AARON MUELLER

---

<b>CONTACT</b>	Khoury College of Computer Sciences Northeastern University 177 Huntington Ave., 22 Fl. Boston, MA 02115 (USA)	<a href="mailto:aa.mueller@northeastern.edu">aa.mueller@northeastern.edu</a> <a href="https://aaronmueller.github.io">aaronmueller.github.io</a> <a href="https://github.com/aaronmueller">github.com/aaronmueller</a>
<b>RESEARCH INTERESTS</b>	<ul style="list-style-type: none"><li>• Natural language processing</li><li>• Robust generalization</li><li>• Mechanistic interpretability</li><li>• Computational psycholinguistics</li></ul>	
<b>EDUCATION</b>	<b>Johns Hopkins University</b> Ph.D., Computer Science M.S.E., Computer Science GPA: 3.9/4.0 <i>Thesis:</i> Emergent Syntactic Behaviors and Mechanisms in Neural Language Models <i>Advisors:</i> Tal Linzen, Mark Dredze	Baltimore, MD Aug. 2023 May 2020
	<b>New York University</b> Visiting academic, Center for Data Science <i>Advisor:</i> Tal Linzen	New York, NY Aug. 2021 – Aug. 2023
	<b>University of Kentucky</b> B.S., Computer Science. <i>Honors</i> B.S., Linguistics. <i>Honors</i> GPA: 4.0/4.0. <i>Summa cum laude</i>	Lexington, KY May 2018 May 2018
<b>ACADEMIC POSITIONS</b>	<b>Northeastern University</b> <i>Zuckerman Postdoctoral Fellow</i> , Khoury College of Computer Sciences <i>Advisor:</i> David Bau	Boston, MA Aug. 2023 – Present
	<b>Technion – Israel Institute of Technology</b> <i>Zuckerman Postdoctoral Fellow</i> , Department of Computer Science <i>Advisor:</i> Yonatan Belinkov	Haifa, Israel Aug. 2023 – Present
<b>INDUSTRY EXPERIENCE</b>	<b>Meta</b> <i>Research Intern</i> <i>Manager:</i> Kanika Narang <ul style="list-style-type: none"><li>– Research in retrieval-augmented generative models for few-shot question answering.</li><li>– Resulted in improved <math>F_1</math> on multiple QA and classification datasets using far fewer parameters than state-of-the-art models. Also resulted in a publication at ACL [3].</li></ul>	Menlo Park, CA May – Nov. 2022
	<b>Amazon Web Services (AWS)</b> <i>Applied Scientist Intern</i> <i>Manager:</i> Saab Mansour <ul style="list-style-type: none"><li>– Research in pre-training methods for improving goal-oriented dialogue agents.</li><li>– Resulted in state-of-the-art few-shot intent classification accuracy (&gt;30% 1-shot gains) and a publication at ACL [8].</li></ul>	Santa Clara, CA May – Aug. 2021
	<b>Raytheon BBN Technologies</b> <i>Research Intern</i> <i>Manager:</i> Ilana Heintz <ul style="list-style-type: none"><li>– Research in low-resource cross-lingual word alignment and entity linking.</li><li>– Implemented convolutional neural machine translation models rivaling our prior seq2seq model’s BLEU with over 20% faster training and over 50% faster inference.</li></ul>	Cambridge, MA May – Aug. 2019

## PUBLICATIONS Peer-reviewed Articles

1. Eric Todd, Millicent L. Li, Arnab Sen Sharma, **Aaron Mueller**, Byron C. Wallace, David Bau. “[Function Vectors in Large Language Models](#).” In *Proceedings of the International Conference on Learning Representations (ICLR)*, 2024.
2. **Aaron Mueller**, Tal Linzen. “[How to Plant Trees in Language Models: Data and Architectural Effects on the Emergence of Syntactic Inductive Biases](#).” In *Proceedings of the Association for Computational Linguistics (ACL)*, 2023.
3. **Aaron Mueller**, Kanika Narang, Lambert Mathias, Qifan Wang, Hamed Firooz. “[Meta-training with Demonstration Retrieval for Efficient Few-shot Learning](#).” In *Findings of the Association for Computational Linguistics (ACL)*, 2023.
4. Koustuv Sinha, Jon Gauthier, **Aaron Mueller**, Kanishka Misra, Keren Fuentes, Roger Levy, Adina Williams. “[Language Model Acceptability Judgements Are Not Always Robust to Context](#).” In *Proceedings of the Association for Computational Linguistics (ACL)*, 2023. **Outstanding Paper Award**.
5. Ian R. McKenzie, Alexander Lyzhov, Michael Martin Pieler, Alicia Parrish, **Aaron Mueller**, Ameya Prabhu, Euan McLean, Xudong Shen, Joe Cavanagh, Andrew George Gritsevskiy, Derik Kauffman, Aaron T. Kirtland, Zhengping Zhou, Yuhui Zhang, Sicong Huang, Daniel Wurgaft, Max Weiss, Alexis Ross, Gabriel Recchia, Alisa Liu, Jiacheng Liu, Tom Tseng, Tomasz Korbak, Najoung Kim, Samuel R. Bowman, Ethan Perez. “[Inverse Scaling: When Bigger Isn’t Better](#).” In *Transactions on Machine Learning Research (TMLR)*, 2023. **Featured Paper**.
6. Julian Michael, Ari Holtzman, Alicia Parrish, **Aaron Mueller**, Alex Wang, Angelica Chen, Divyam Madaan, Nikita Nangia, Richard Yuanzhe Pang, Jason Phang, Samuel R. Bowman. “[What Do NLP Researchers Believe? Results of the NLP Community Metasurvey](#).” In *Proceedings of the Association for Computational Linguistics (ACL)*, 2023.
7. **Aaron Mueller**, Robert Frank, Tal Linzen, Luheng Wang, Sebastian Schuster. “[Coloring the Blank Slate: Pre-training Imparts a Hierarchical Inductive Bias to Sequence-to-sequence Models](#).” In *Findings of the Association for Computational Linguistics (ACL)*, 2022.
8. **Aaron Mueller**, Jason Krone, Salvatore Romeo, Saab Mansour, Elman Mansimov, Yi Zhang, Dan Roth. “[Label Semantic Aware Pre-training for Few-shot Text Classification](#).” In *Proceedings of the Association for Computational Linguistics (ACL)*, 2022.
9. **Aaron Mueller**, Yu Xia, Tal Linzen. “[Causal Analysis of Syntactic Agreement Neurons in Multilingual Language Models](#).” In *Proceedings of the Conference on Computational Natural Language Learning (CoNLL)*, 2022.
10. Alexandra DeLucia, Shijie Wu, **Aaron Mueller**, Carlos Aguirre, Mark Dredze, Philip Resnik. “[BERNICE: A Multilingual Pre-trained Encoder for Twitter](#).” In *Proceedings of the Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2022.
11. **Aaron Mueller**, Mark Dredze. “[Fine-tuning Encoders for Improved Monolingual and Zero-shot Polylingual Neural Topic Modeling](#).” In *Proceedings of the North American Chapter of the Association for Computational Linguistics (NAACL)*, 2021.
12. **Aaron Mueller**, Zach Wood-Doughty, Silvio Amir, Mark Dredze, Alicia L. Nobles. “[Demographic Representation and Collective Storytelling in the Me Too Twitter Hashtag Activism Movement](#).” In *Proceedings of the Association for Computing Machinery (ACM) on Human-Computer Interaction (HCI)*, vol. CSCWI, 2021.
13. Matthew Finlayson\*, **Aaron Mueller\***, Sebastian Gehrmann, Stuart Shieber, Tal Linzen, Yonatan Belinkov. “[Causal Analysis of Syntactic Agreement Mechanisms in Neural Language Models](#).” In *Proceedings of the Association for Computational Linguistics (ACL)*, 2021. [\*Equal contribution]
14. Alexandra DeLucia\*, **Aaron Mueller\***, Xiang Lisa Li, João Sedoc. “[Decoding Methods for](#)

[Neural Narrative Generation.](#)” In *Proceedings of the Workshop on Generation Evaluation and Metrics (GEM) at Association for Computational Linguistics (ACL)*, 2021. [\*Equal contribution]

15. **Aaron Mueller**, Garrett Nicolai, Panayiota Petrou-Zeniou, Natalia Talmina, Tal Linzen. “[Cross-linguistic Syntactic Evaluation of Word Prediction Models.](#)” In *Proceedings of the Association for Computational Linguistics (ACL)*, 2020.
16. **Aaron Mueller**, Garrett Nicolai, Arya D. McCarthy, Dylan Lewis, Winston Wu, David Yarowsky. “[An Analysis of Massively Multilingual Neural Machine Translation for Low-Resource Languages.](#)” In *Proceedings of the Language Resources and Evaluation Conference (LREC)*, 2020.
17. Arya D. McCarthy, Rachel Wicks, Dylan Lewis, **Aaron Mueller**, Winston Wu, Oliver Adams, Garrett Nicolai, Matt Post, David Yarowsky. “[The Johns Hopkins University Bible Corpus: 1600+ Tongues for Typological Exploration.](#)” In *Proceedings of the Language Resources and Evaluation Conference (LREC)*, 2020.
18. Garrett Nicolai, Dylan Lewis, Arya D. McCarthy, **Aaron Mueller**, Winston Wu, David Yarowsky. “[Fine-grained Morphosyntactic Analysis and Generation Tools for More Than One Thousand Languages.](#)” In *Proceedings of the Language Resources and Evaluation Conference (LREC)*, 2020.
19. Marten van Schijndel, **Aaron Mueller**, Tal Linzen. “[Quantity Doesn’t Buy Quality Syntax with Neural Language Models.](#)” In *Proceedings of the Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2019.
20. Arya D. McCarthy, Winston Wu, **Aaron Mueller**, Bill Watson, David Yarowsky. “[Modeling Color Terminology Across Thousands of Languages.](#)” In *Proceedings of the Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2019.
21. **Aaron Mueller\***, Yash Kumar Lal\*. “[Sentence-Level Adaptation for Low-Resource Neural Machine Translation.](#)” In *Proceedings of the Workshop on Technologies for Machine Translation of Low-Resource Languages (LoResMT) at Machine Translation Summit (MTSummit)*, 2019. [\*Equal contribution]

#### Invited Publications

22. Alex Warstadt\*, **Aaron Mueller\***, Leshem Choshen, Ethan Wilcox, Juan Ciro, Rafael Mosquera, Bhargavi Paranjabe, Adina Williams, Tal Linzen, Ryan Cotterell. “[Findings of the BabyLM Challenge: Sample-efficient Pretraining on Developmentally Plausible Corpora.](#)” Shared task proceedings in *Conference on Computational Natural Language Learning (CoNLL)*, 2023. [\*Equal contribution]

#### Preprints & In Submission

23. **Aaron Mueller**, Albert Webson, Jackson Petty, Tal Linzen. “[In-context Learning Generalizes, But Not Always Robustly: The Case of Syntax.](#)” In submission to *North American Chapter of the Association for Computational Linguistics (NAACL)*, 2024.

#### Other Articles

24. **Aaron Mueller**. “Emergent Syntactic Behaviors and Mechanisms in Neural Language Models.” Ph.D. Dissertation, Johns Hopkins University, 2023. Committee: Tal Linzen, Mark Dredze, David Yarowsky, Yonatan Belinkov.

#### INVITED TALKS

*Evaluating and Surgically Improving Generalization in Language Models.*  
NLP Seminar, University of Pittsburgh (Pittsburgh, PA). Feb. 29, 2024.

*Evaluating and Surgically Improving Generalization in Language Models.*  
Deep Learning Superlab, Brown University (Providence, RI). Feb. 15, 2024.

*Planting Trees in Language Models: Emergent Syntactic Behaviors and Mechanisms from Pre-training.*  
Koller Lab, Saarland University (Saarbrücken, Germany). Feb. 7, 2023.

*Planting Trees in Language Models: Emergent Syntactic Behaviors and Mechanisms from Pre-training.*

NLP Seminar, Technion – Israel Institute of Technology (Haifa, Israel). Dec. 14, 2022.

*Planting Trees in Language Models: Emergent Syntactic Behaviors and Mechanisms from Pre-training.*

Bar-Ilan NLP Seminar, Bar-Ilan University (Ramat Gan, Israel). Dec. 13, 2022.

*What Generalizations do Sequence-to-sequence Models Learn from Multilingual Text? Insights from Translation and Syntactic Transformations.*

Multilingual Text Processing Group, National Research Council of Canada (Ottawa, ON). Mar. 4, 2022.

*Syntactic Agreement in Neural Language Models: How Well and Where Do They Perform Subject-Verb Agreement?*

Language & Understanding Group, Mila – Québec Artificial Intelligence Institute (Montréal, QC). Mar. 22, 2021.

*Causal Mediation Analysis for Analyzing Neural Networks.*

Fairness & Interpretability Research Talk Series, Google (New York, NY). Mar. 17, 2021.

*Causal Analysis of Syntactic Agreement Mechanisms in Neural Language Models.*

Center for Language & Speech Processing Seminar, Johns Hopkins University (Baltimore, MD). Feb. 12, 2021.

<b>FELLOWSHIPS AND AWARDS</b>	<b>Zuckerman Fellow, International</b>	2023–2025
	Two-year postdoctoral fellowship. Supports research joint with an Israeli university and an American university. (\$126,000)	
	<b>Microsoft Accelerate Foundation Models Research Award, International</b>	2023
	Awarded for research on the capabilities of large language models. Provides OpenAI API credits and priority GPT-4 access. (\$10,000)	
	<b>National Science Foundation Graduate Research Fellow, National</b>	2018 - 2023
	Five-year graduate research fellowship. Provides three years of Ph.D. funding. (\$135,000)	
	<b>Gaines Fellow, University of Kentucky</b>	2016 - 2018
	Two-year fellowship. Requires the completion of a juried project, a thesis project, and a seminar in the humanities. (\$5,000)	
	<b>Patterson Scholar, University of Kentucky</b>	2014 - 2018
	Four-year scholarship covering tuition, educational materials, and room & board. Awarded to undergraduates who have earned National Merit semifinalist standing or higher. (\$86,000)	
	<b>Goldwater Scholarship (Honorable Mention), National</b>	2017
	<b>Phi Beta Kappa, National</b>	2017
	<b>Raymond F. Betts Scholar, University of Kentucky</b>	2017
	Awarded for thesis research. Used funds to design language technologies for low-resource dialects of French. (\$2,500)	
	<b>Linguistics Research Award, University of Kentucky</b>	2016
<b>MENTORING</b>	Awarded to an undergraduate to facilitate a year-long research project in linguistics. (\$500)	
	Ph.D. students:	
	– Aruna Sankaranarayanan (MIT). Joint with Forrest Davis. Research on natural and artificial grammar learning in language models.	2023–Present

	<ul style="list-style-type: none"> <li>– Eric W. Todd (Northeastern). Joint with David Bau. Research on how functions are represented in neural language models. Resulted in a publication at ICLR [1].</li> <li>– Juan Diego Rodriguez (UT Austin). Joint with Kanishka Misra. Research in how concepts are organized in language models.</li> </ul>	<p>2023–Present</p> <p>2023–Present</p>
	Master’s students:	
	<ul style="list-style-type: none"> <li>– Dan Pechi (NYU). Research on imparting better inductive biases to language models.</li> <li>– Swapnil Sharma (NYU). Research on evaluating summarization models.</li> <li>– Yash Kumar Lal (JHU). Resulted in a workshop publication [21].</li> </ul>	<p>2023</p> <p>2022–2023</p> <p>2018–2019</p>
	Undergraduate researchers:	
	<ul style="list-style-type: none"> <li>– Yu Xia (NYU). Resulted in a publication at CoNLL [9].</li> <li>– Matthew Finlayson (Harvard). Resulted in a publication at ACL [13].</li> </ul>	<p>2021–2022</p> <p>2020–2021</p>
<b>TEACHING</b>	<b>Johns Hopkins University</b> <i>Teaching Assistant</i> <i>Instructor:</i> Mathias Unberath <ul style="list-style-type: none"> <li>– Machine Learning: AI System Design &amp; Development</li> </ul> <b>New York University</b> <i>Guest Lecture</i> <i>Instructor:</i> Tal Linzen <ul style="list-style-type: none"> <li>– Computational Linguistics &amp; Cognitive Science</li> </ul>	<p>Spring 2020</p> <p>Spring 2023</p>
<b>SERVICE</b>	Organizing Committees: <ul style="list-style-type: none"> <li>– <a href="#">The BabyLM Challenge</a> (2023, 2024)</li> <li>– <a href="#">The Inverse Scaling Prize</a> (2022)</li> </ul> Reviewing: <ul style="list-style-type: none"> <li>– ACL Rolling Review (Oct. 2021 – present; monthly)</li> <li>– NAACL (2024, 2021)</li> <li>– ACL (2023, 2022, 2020)</li> <li>– EMNLP (2023, 2022, 2019)</li> <li>– CoNLL (2023, 2022)</li> <li>– FAccT (2024)</li> <li>– TACL (2022)</li> <li>– CSCW (2021)</li> <li>– COLING (2020)</li> </ul>	
<b>SKILLS</b>	Programming: <ul style="list-style-type: none"> <li>– Languages: Python, C++, HTML, CSS, Javascript, Bash</li> <li>– Machine Learning Toolkits: PyTorch (incl. HuggingFace, fairseq, sockeye), NLTK, Scikit-learn, numpy</li> <li>– Version Control: DVCS (Git, Bitbucket)</li> </ul> Linguistic Tools: <ul style="list-style-type: none"> <li>– Praat, AntConc, QGIS, Audacity</li> </ul>	
<b>NATURAL LANGUAGES</b>	English (native language). French (B2, Canadian).	